

# Boost Apache Kafka with Schema Registry

Apache Kafka  
doesn't care

```

{
  "id": 778,
  "shop": "Luigis Pizza",
  "name": "Edward Olson",
  "phoneNumbers":
    ["(935)503-3765x4154","(935)12345"],
  "address": "Unit 9398 Box 2056\nDP0 AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍖", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍌", "🌶️", "🍍"]
    }
  ]
}

```

Carlo, Margherita, 2  
 Francesco, Salami, 1

dGVzdCx0ZXN0MSx0ZXN0Mix0ZXN0LHRlc3QxLHRlc3Q  
 yLHRlc3QsdGVzdDEsdGVzdDI

# The Beginning

# JSON

```
{
  "id": 778,
  "shop": "Luigis Pizza",
  "name": "Edward Olson",
  "phoneNumbers":
    ["(935)503-3765x4154","(935)12345"],
  "address": "Unit 9398 Box 2056\nDPO AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍗", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍌", "🌶️", "🍍"]
    }
  ]
}
```

A bit  
Heavy?

# JSON

```
{
  "id": 778,
  "shop": "Luigis Pizza",
  "name": "Edward Olson",
  "phoneNumbers":
    ["(935)503-3765x4154","(935)12345"],
  "address": "Unit 9398 Box 2056\nDPO AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍖", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍌", "🌶️", "🍍"]
    }
  ]
}
```

# Schema

```
{
  "id": int,
  "shop": string,
  "name": string,
  "phoneNumbers":
    [string],
  "address": string,
  "image": string,
  "pizzas": [
    {
      "pizzaName": string,
      "additionalToppings": [string]
    }
  ]
}
```

# Payload

```
{
  778,
  "Luigis Pizza",
  "Edward Olson",
  ["(935)503-3765x4154","(935)12345"],
  "Unit 9398 Box 2056\nDPO AP 24022",
  null,
  [
    {
      "Salami",
      ["🍖", "🌶️"]
    },
    {
      "Margherita",
      ["🍌", "🌶️", "🍍"]
    }
  ]
}
```



```
{
  "id": 778,
  "shop": "Luigis Pizza",
  "name": "Edward Olson",
  "phoneNumbers":
    ["(935)503-3765x4154","(935)12345"],
  "address": "Unit 9398 Box 2056\nDP0 AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍖", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍌", "🌶️", "🍍"]
    }
  ]
}
```



```
{
  778,
  "Luigis Pizza",
  "Edward Olson",
  ["(935)503-3765x4154","(935)12345"],
  "Unit 9398 Box 2056\nDP0 AP 24022",
  null,
  [
    {
      "Salami",
      ["🍖", "🌶️"]
    },
    {
      "Margherita",
      ["🍌", "🌶️", "🍍"]
    }
  ]
}
```

# Benefits

Less data in transit

Less data stored

Network is a bottleneck

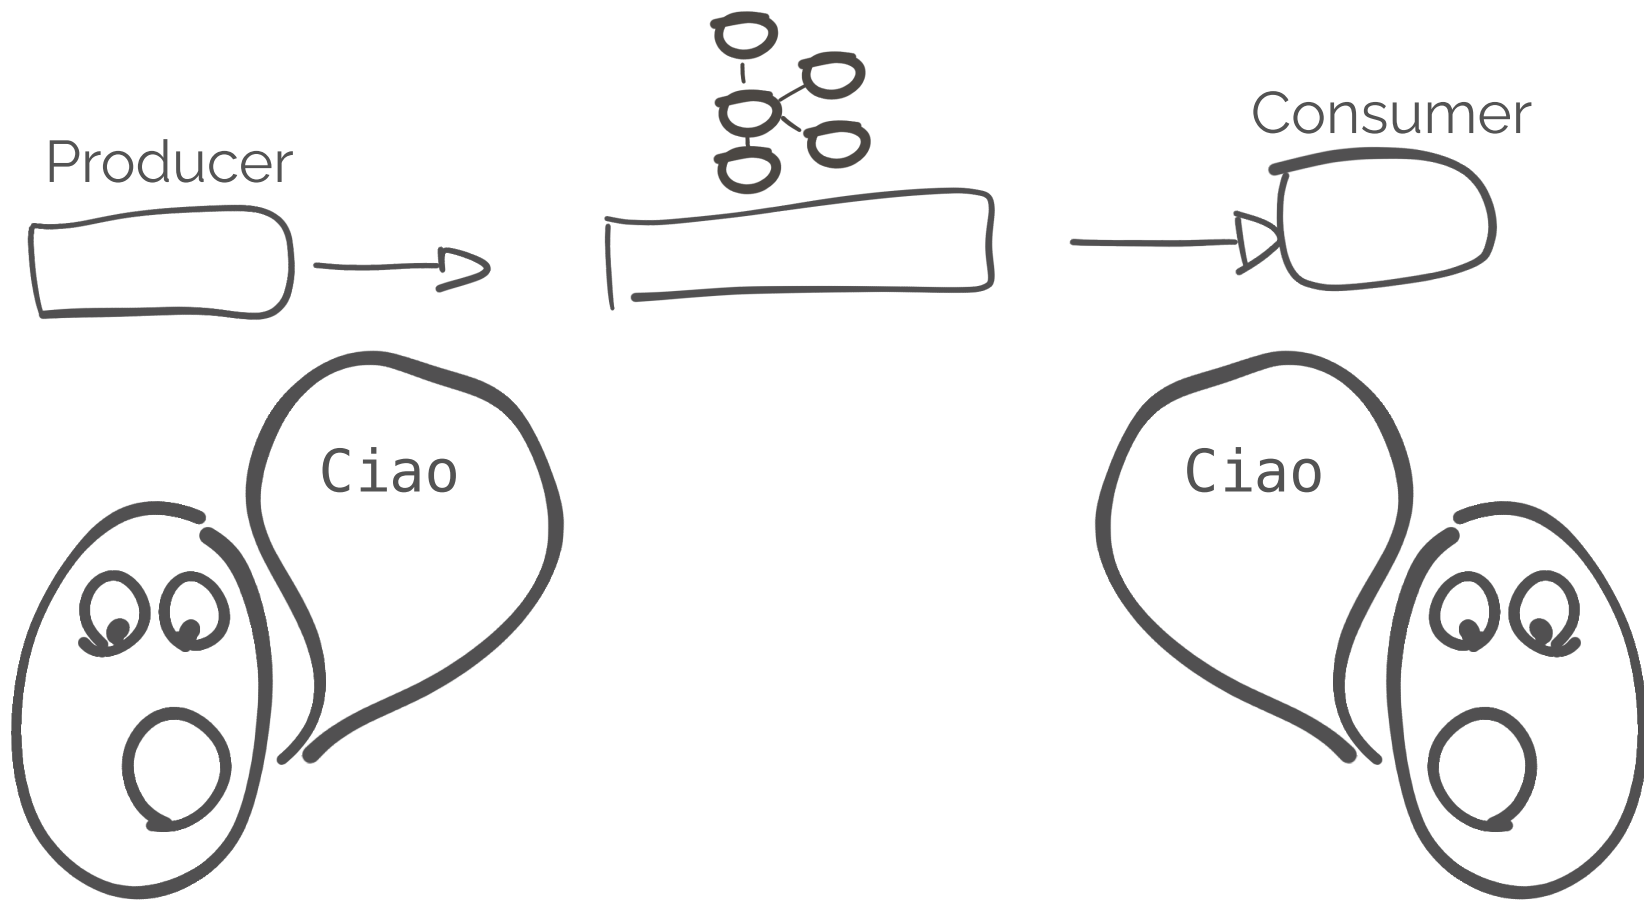
better performance

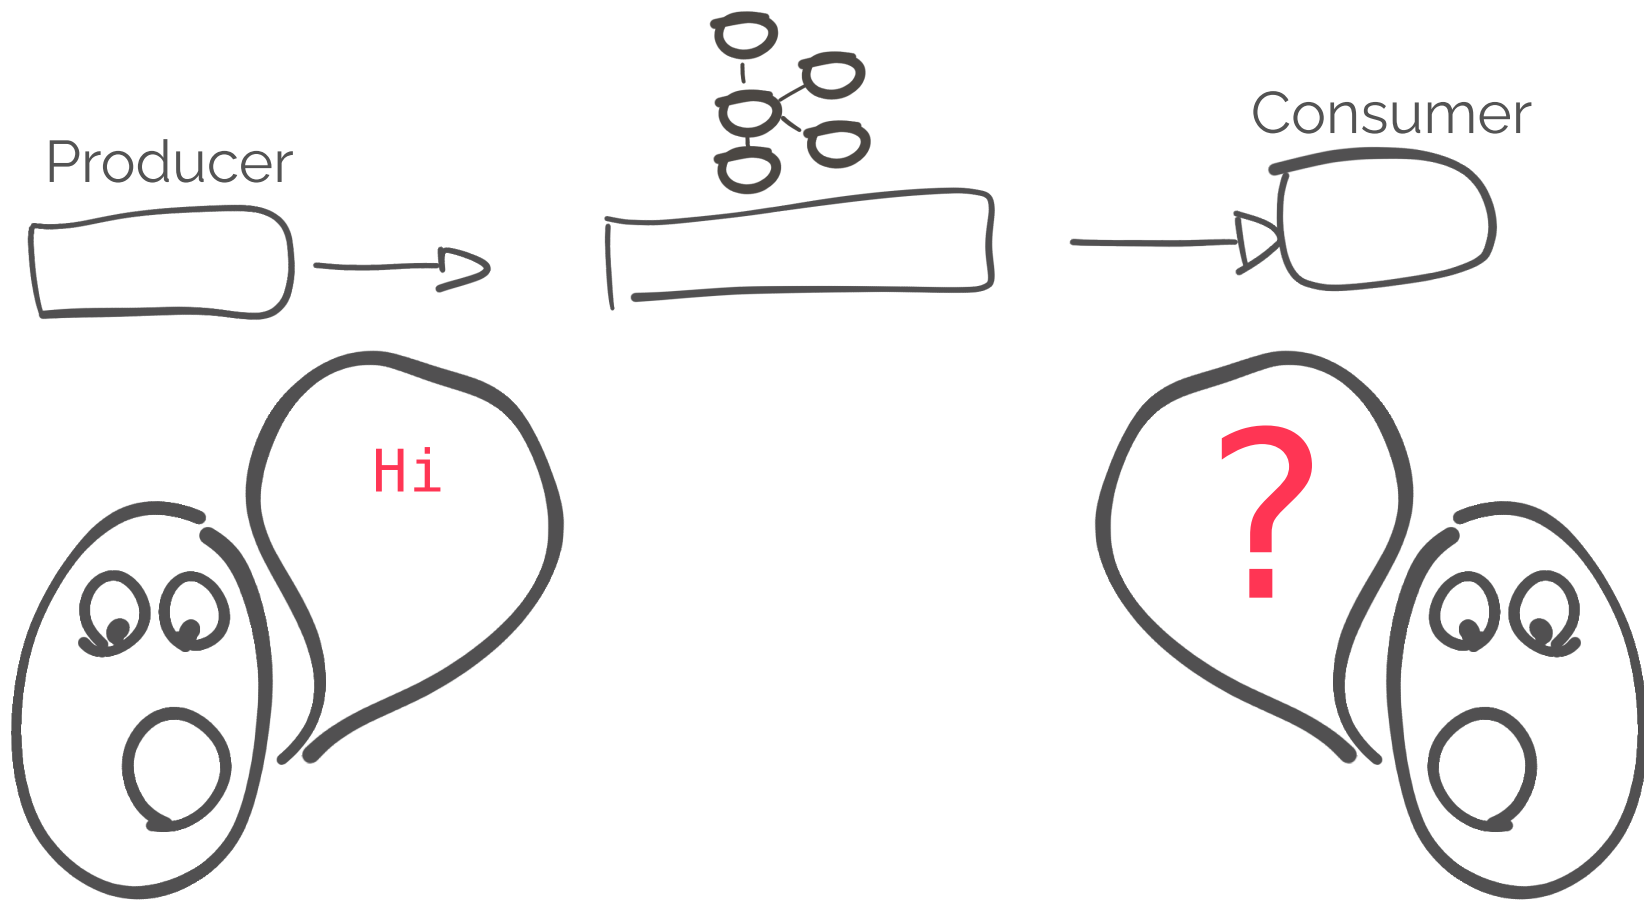
# Cons

You need to translate

Is it only about  
**Compression?**

**No**  
**input**  
**validation**





# JSON

```
{
  "id": 778,
  "shop": "Luigis Pizza",
  "name": "Edward Olson",
  "phoneNumbers":
    ["(935)503-3765x4154","(935)12345"],
  "address": "Unit 9398 Box 2056\nDPO AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍗", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍌", "🌶️", "🍍"]
    }
  ]
}
```

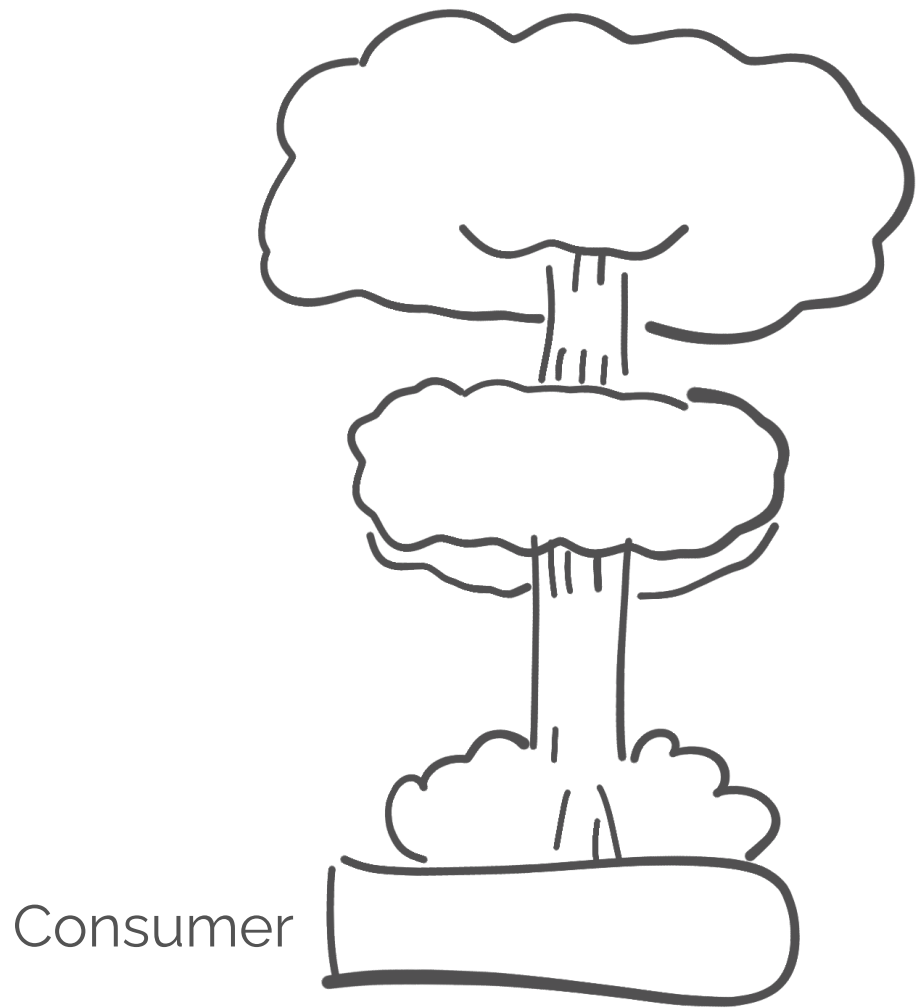


# JSON

```
{
  "id": 778,
  "shop": "Luigis Pizza",
  "name": "Edward Olson",
  "phoneNumbers":
    ["(935)503-3765x4154","(935)12345"],
  "address": "Unit 9398 Box 2056\nDPO AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍖", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍕", "🌶️", "🍍"]
    }
  ]
}
```

# JSON

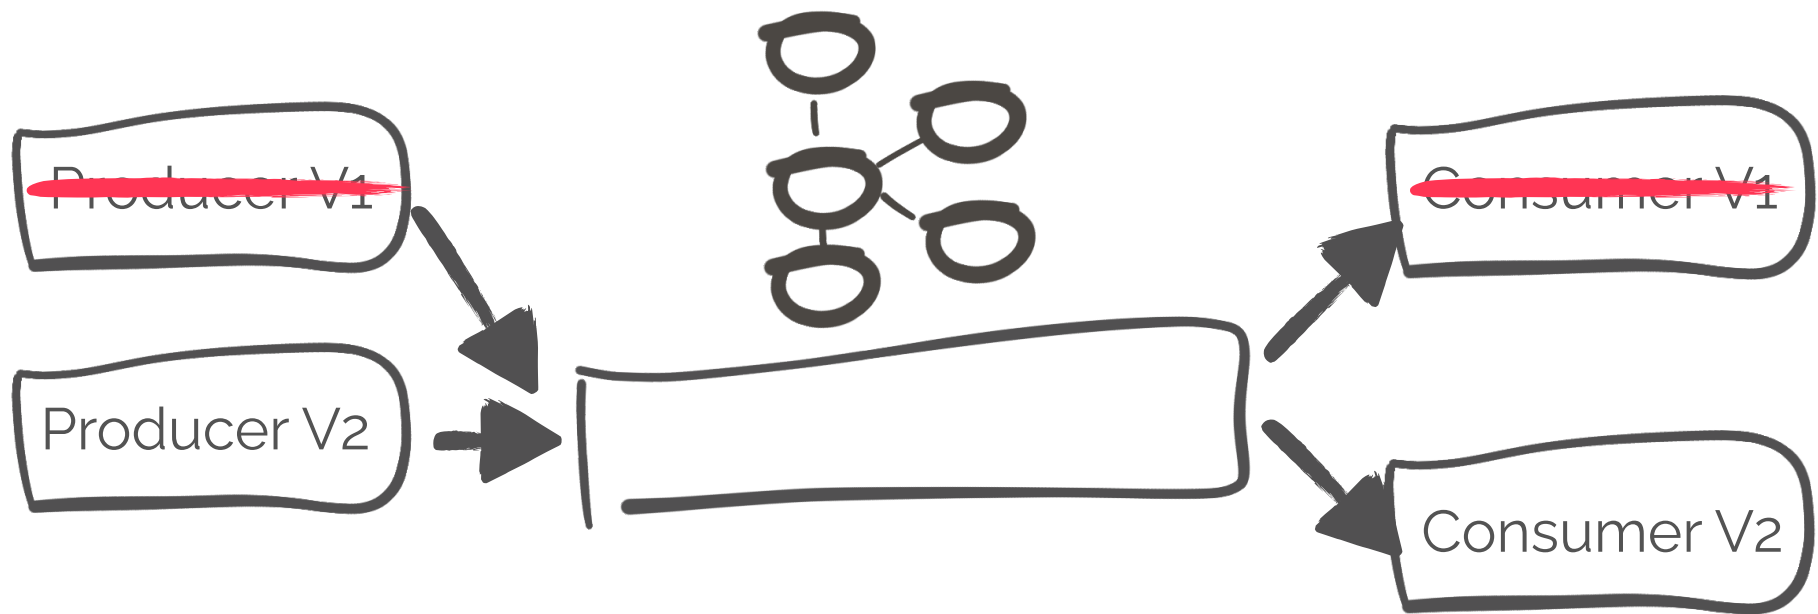
```
{
  "id": 778,
  "shop": "Luigis Pizza",
  "firstname": "Edward",
  "lastname": "Olson",
  "phoneNumbers":
    ["(935)503-3765x4154", "(935)12345"],
  "address": "Unit 9398 Box 2056\nDPO AP 24022",
  "image": null,
  "pizzas": [
    {
      "pizzaName": "Salami",
      "additionalToppings": ["🍌", "🌶️"]
    },
    {
      "pizzaName": "Margherita",
      "additionalToppings": ["🍌", "🌶️", "🍍"]
    }
  ]
}
```



# Tight Coupling

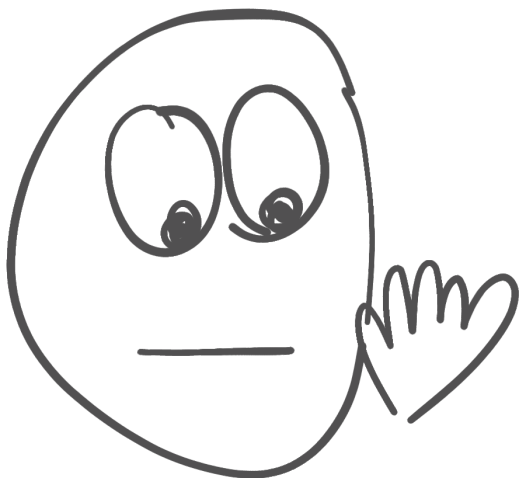
# Demo 1

# Evolution





Allowed Change



Forbidden Change



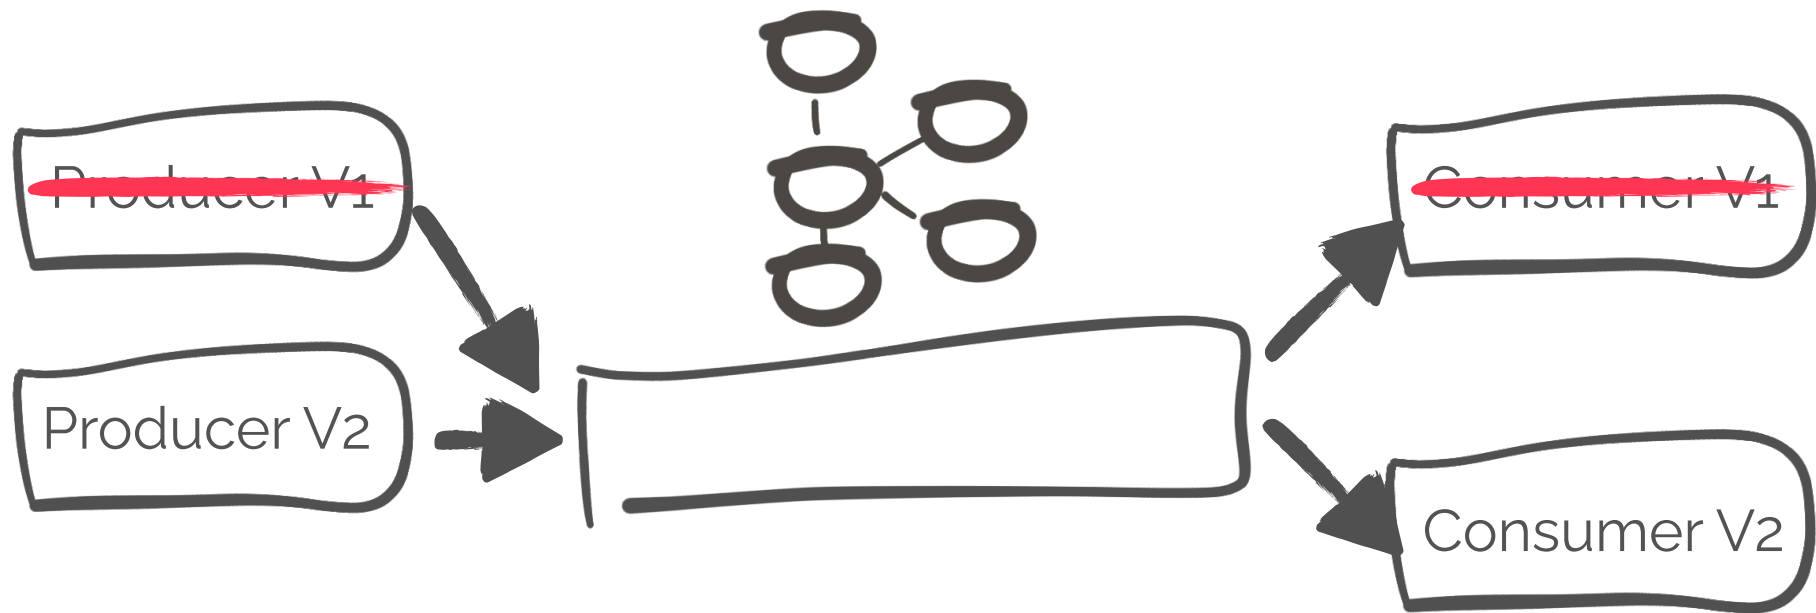
# Schema Registry

# Topic A

Schema

Compatibility

# Compatibility



COMPATIBILITY = NONE



$V_{x-2}$     $V_{x-1}$     $V_x$     $V_{x+1}$     $V_{x+2}$

No checks → Like JSON

COMPATIBILITY = BACKWARD



## Allowed Changes

- Delete Fields
- Add Optional Fields

Consumers using the V4 schema to decode the message, will be able to also parse messages sent with V3 schema

COMPATIBILITY = BACKWARD\_TRANSITIVE



$V_2$

$V_3$

$V_4$

$V_5$

$V_6$

## Allowed Changes

- Delete Fields
- Add Optional Fields

Consumers using the V4 schema to decode the message, will be able to also parse messages sent with V3, V2, V1 schema

COMPATIBILITY = FORWARD

V<sub>2</sub>

V<sub>3</sub>

V<sub>4</sub>

V<sub>5</sub>

V<sub>6</sub>

## Allowed Changes

- Add Fields
- Delete Optional Fields

Consumers using V<sub>4</sub> schema will be able to also parse messages produced with V<sub>5</sub> schemas

COMPATIBILITY = FORWARD\_TRANSITIVE

V<sub>2</sub>

V<sub>3</sub>

V<sub>4</sub>

V<sub>5</sub>

V<sub>6</sub>

## Allowed Changes

- Add Fields
- Delete Optional Fields

Consumers using V<sub>4</sub> schema will be able to also parse messages produced with V<sub>5</sub>, V<sub>6</sub>, ... schemas [@ftisiot](#) | [@aiven\\_io](#)



COMPATIBILITY = FULL

$V_2$

$V_3$

$V_4$

$V_5$

$V_6$

## Allowed Changes

- Add Optional Fields
- Delete Optional Fields

COMPATIBILITY = FULL\_TRANSITIVE



$V_2$

$V_3$

$V_4$

$V_5$

$V_6$

## Allowed Changes

- Add Optional Fields
- Delete Optional Fields

# Compatibility Summary

## BACKWARD

"The consumer i just started, will be able to read all the history"

If you want to be sure to read the past

## FORWARD

The consumer i just started, will be able to read all the future

If you don't want to break previous consumers

# Schema Registry Summary

Detach Schema from Data

Allow Evolution

Define Compatibility

# Schemas... Additional Benefit

Work OOTB with Kafka Connect

# Karapace

Schema Registry

REST APIs

Fully Open Source

<https://www.karapace.io/>

@ftisiot | @aiven\_io

# Karapace

AVRO

JSON

Protobuf

```

from confluent_kafka import Producer
from confluent_kafka.serialization import StringSerializer, SerializationContext, MessageField
from confluent_kafka.schema_registry import SchemaRegistryClient
from confluent_kafka.schema_registry.avro import AvroSerializer

value_str = """
{
  "namespace": "example.avro",
  "type": "record",
  "name": "User",
  "fields": [
    {"name": "name", "type": "string"},
    {"name": "age", "type": ["int", "null"]},
    {"name": "nationality", "type": ["string", "null"], "default": "Italian"}
  ]
}
"""

```

```

def user_to_dict(user, ctx):
    return dict(
        name=user.name, age=user.age, nationality=user.nationality
    )

```

```

# Create the Avro serializer
avro_serializer_value = AvroSerializer(
    schema_registry_client, value_str, user_to_dict
)

```

```

payload = User(name="John", age=30)
topic = "test-avro"
producer.produce(
    topic=topic,
    value=avro_serializer_value(
        payload,
        SerializationContext(topic, MessageField.VALUE),
    ),
)

```

# Producer



```
# Import required libraries
from confluent_kafka import DeserializingConsumer
from confluent_kafka.serialization import SerializationContext, MessageField, StringSerializer
from confluent_kafka.schema_registry import SchemaRegistryClient
from confluent_kafka.schema_registry.avro import AvroDeserializer

avro_deserializer_value = AvroDeserializer(
    schema_registry_client, value_str, dict_to_user
)

# Subscribe to the Kafka topic
consumer.subscribe([env.TOPIC_PREFIX + '-avro'])

# Consume messages from the Kafka topic
while True:
    try:
        msg = consumer.poll(1.0)

        if msg is None:
            continue

        user = avro_deserializer_value(
            msg.value(), SerializationContext(msg.topic(), MessageField.VALUE)
        )

        if user is not None and key is not None:
            print("Key --> {}\n User record --> name: {} age: {}\n".format(key.id, user.name, user.age))
    except KeyboardInterrupt:
        break
consumer.commit()
```

# Consumer

# How to Kafka Connect

```
"value.converter": "io.confluent.connect.avro.AvroConverter",  
"value.converter.schema.registry.url": "[HOST]:[PORT]",  
"value.converter.basic.auth.credentials.source": "USER_INFO",  
"value.converter.schema.registry.basic.auth.user.info": "[USERNAME]:[PASSWORD]"
```

# Need to enhance throughput?

Keys

Batching

# Enhance Apache Kafka with Schemas