



# Apache Kafka and Flink


## Stateful Streaming Data Pipelines made easy with SQL


PRESS ME!












**What  
is  
Apache  
Kafka?**



Topic A




Topic B









# kafka-python

```
from kafka import KafkaProducer
```

```
producer = KafkaProducer(  
    bootstrap_servers=['broker1:1234']  
)
```

```
producer.send(  
    'my-topic-name',  
    b'my-message'  
)
```

```
producer.flush()
```

```
{
```


```
  "id": 1,  
  "shop": "Mario's Pizza",  
  "name": "Arsenio Pisaroni-Boccaccio",  
  "phoneNumber": "+39 51 0290746",  
  "address": "Via Ugo 01, Montegrotto, 85639 Padova(PD)",  
  "pizzas": [  
    {  
      "pizzaName": "Margherita",  
      "additionalToppings": ["ham"]  
    },  
    {  
      "pizzaName": "Diavola",  
      "additionalToppings": ["mozzarella", "banana", "onion"]  
    }  
  ]  
}
```







# Create Stateful Apps



Kafka Streams






Faust




KSQL



Apache Flink









SQL





Table API



DataStream API



Filter



Join



Aggregate




Explode




Detect





Change  
Shape





# Connect Flink



```
{  
  "id": 1,  
  "shop": "Mario's Pizza",  
  "name": "Arsenio Pisaroni-Boccaccio",  
  "phoneNumber": "+39 51 0290746",  
  "address": "Via Ugo 01, Montegrotto, 85639 Padova(PD)",  
  "pizzas": [  
    {  
      "pizzaName": "Margherita",  
      "additionalToppings": ["ham"]  
    }  
  ]  
}
```

pizza_name	base_price
Marinara	4
Diavola	6
Mari & Monti	8
Salami	7
Peperoni	8
Margherita	5

# Kafka Source

```
CREATE TABLE pizza_orders (  
  id INT,  
  shop VARCHAR,  
  name VARCHAR,  
  phoneNumber VARCHAR,  
  address VARCHAR,  
  pizzas ARRAY  
    <ROW (  
      pizzaName VARCHAR,  
      additionalToppings ARRAY <VARCHAR>)>  
) WITH (  
  'connector' = 'kafka',  
  'properties.bootstrap.servers' = 'kafka:13041',  
  'topic' = 'pizza-orders',  
  'scan.startup.mode' = 'earliest-offset',  
  ...  
);
```

# Pg Source

```
CREATE TEMPORARY TABLE pizza_prices (  
  pizza_name VARCHAR,  
  base_price INT,  
  PRIMARY KEY (pizza_name) NOT ENFORCED  
) WITH (  
  'connector' = 'jdbc',  
  'url' = 'jdbc:postgresql://pghost:13039/db',  
  'username'='avnadmin',  
  'password'='verysecurepassword123',  
  'table-name' = 'pizza_price'  
);
```

# Pg Tgt


```
CREATE TABLE order_price (  
  id INT,  
  pizza_name VARCHAR,  
  base_price INT,  
  nr_pizzas BIGINT NOT NULL,  
  PRIMARY KEY (id, pizza_name) NOT ENFORCED  
) WITH (  
  'connector' = 'jdbc',  
  'url' = 'jdbc:postgresql://pghost:13039/db',  
  'username'='avnadmin',  
  'password'='verysecurepassword123',  
  'table-name' = 'order_price'  
);
```




# Create Pipeline

```
insert into order_price
select id,
       b.pizzaName,
       base_price,
       count(*) nr_pizzas
from pizza_orders cross join UNNEST(pizzas) b
LEFT OUTER JOIN pizza_prices
  FOR SYSTEM_TIME AS OF orderProctime AS pp
  ON b.pizzaName = pp.pizza_name
group by id,
       b.pizzaName,
       base_price;
```





# Resources



<https://kafka.apache.org/>



<http://flink.apache.org/>



<https://github.com/aiven/flink-sql-cli-docker>



<https://aiven.io/blog/build-a-streaming-sql-pipeline-with-flink-and-kafka>



<https://aiven.io>